

# StudentsExamsScoreAnalysis

November 20, 2025

The program performs exploratory data analysis and predictive modeling on a dataset of students' exam scores. It loads the data from student\_exam\_scores.csv into a pandas DataFrame, creates a derived categorical variable isPassed, and computes descriptive statistics for the main features. Relationships between study habits and performance are explored using scatter plots with a pass/fail hue, linear regression plots, correlation heatmaps, and histograms of key variables. The data is then split into training and test sets, and two regression models (Linear Regression and Random Forest) are trained to predict exam\_score. Their performance is evaluated using MSE and R<sup>2</sup> on both train and test sets, the results are summarized in a comparison table, and an additional scatter plot compares predicted vs. actual exam scores for the linear model.

Importing libraries.

```
[28]: import pandas as pd
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.ensemble import RandomForestRegressor
```

Loading data.

```
[29]: scores = pd.read_csv("student_exam_scores.csv", delimiter=",")
```

	student_id	hours_studied	sleep_hours	attendance_percent	previous_scores	exam_score
0	S001	8.0	8.8	72.1	45	30.2
1	S002	1.3	8.6	60.7	55	25.0
2	S003	4.0	8.2	73.7	86	35.8
3	S004	3.5	4.8	95.1	66	34.0
4	S005	9.1	6.4	89.8	71	40.3

Printing information about dataframe.

```
[54]: scores.dtypes
```

```
[54]: student_id          object
hours_studied      float64
sleep_hours        float64
attendance_percent float64
previous_scores    int64
exam_score         float64
isPassed           object
dtype: object
```

```
[55]: scores.head()
```

```
[55]:   student_id  hours_studied  sleep_hours  attendance_percent  previous_scores \
0          S001          8.0          8.8            72.1                  45
1          S002          1.3          8.6            60.7                  55
2          S003          4.0          8.2            73.7                  86
3          S004          3.5          4.8            95.1                  66
4          S005          9.1          6.4            89.8                  71

      exam_score  isPassed
0          30.2    Passed
1          25.0   Failed
2          35.8    Passed
3          34.0    Passed
4          40.3    Passed
```

Printing basic statistics of dataset.

```
[30]: scores.describe()
```

```
[30]:   hours_studied  sleep_hours  attendance_percent  previous_scores \
count      200.000000  200.000000      200.000000  200.000000
mean       6.325500    6.622000      74.830000   66.800000
std        3.227317    1.497138     14.249905   15.663869
min        1.000000    4.000000      50.300000   40.000000
25%        3.500000    5.300000      62.200000   54.000000
50%        6.150000    6.700000      75.250000   67.500000
75%        9.000000    8.025000      87.425000   80.000000
max       12.000000    9.000000     100.000000  95.000000

      exam_score
count  200.000000
mean   33.955000
std    6.789548
min   17.100000
```

```
25%      29.500000
50%      34.050000
75%      38.750000
max      51.300000
```

Checking dataset for null values and dropping duplicates.

```
[32]: scores.isna().sum()
scores.drop_duplicates()
```

```
[32]:    student_id  hours_studied  sleep_hours  attendance_percent \
0          S001           8.0         8.8            72.1
1          S002           1.3         8.6            60.7
2          S003           4.0         8.2            73.7
3          S004           3.5         4.8            95.1
4          S005           9.1         6.4            89.8
..          ...
195         S196          10.5         5.4            94.0
196         S197           7.1         6.1            85.1
197         S198           1.6         6.9            63.8
198         S199          12.0         7.3            50.5
199         S200          10.2         6.3            97.4

      previous_scores  exam_score
0                  45       30.2
1                  55       25.0
2                  86       35.8
3                  66       34.0
4                  71       40.3
..                  ...
195                 87       42.7
196                 92       40.4
197                 76       28.2
198                 58       42.0
199                 68       37.8
```

[200 rows x 6 columns]


Creating a new categorical column “isPassed”.

```
[56]: scores["isPassed"] = np.where(scores["exam_score"] > 30.0, "Passed", "Failed")
```

Visualization of the relationship between hours studied and exam score.

```
[57]: plot = sns.scatterplot(scores, x="hours_studied", y="exam_score",  $\hookrightarrow$ 
                           hue="isPassed")
plot.set(xlabel = "Hours Studied", ylabel = "Exam Score")
```

```
[57]: [Text(0.5, 0, 'Hours Studied'), Text(0, 0.5, 'Exam Score')]
```




Visualization of the relationship between hours studied and exam score using linear regression plot.

```
[58]: plt.figure(figsize = (6,4))
plot = sns.lmplot(scores, x="hours_studied", y="exam_score", line_kws={"color": "orange"})
plot.set(xlabel = "Hours Studied", ylabel = "Exam Score", title="Exam Score vs Hours Studied")
```

```
[58]: <seaborn.axisgrid.FacetGrid at 0x1bcc2abb9d0>
```

```
<Figure size 600x400 with 0 Axes>
```




Visualization of the relationship between attendance percent and exam score using linear regression plot.

```
[35]: plt.figure(figsize = (6,4))
plot = sns.lmplot(scores, x="attendance_percent", y="exam_score", line_kws={"color": "orange"})
plot.set(xlabel = "Attendance Percent", ylabel = "Exam Score", title="Exam Score vs Attendance Percent")
```

```
[35]: <seaborn.axisgrid.FacetGrid at 0x1bcc16e4690>
```

```
<Figure size 600x400 with 0 Axes>
```



Creating a correlation matrix out of numeric values.

```
[61]: corr = scores[["hours_studied", "sleep_hours", "exam_score", "attendance_percent"]].corr()
print(corr)
```

	hours_studied	sleep_hours	exam_score	attendance_percent
hours_studied	1.000000	0.077864	0.776751	-0.031311
sleep_hours	0.077864	1.000000	0.188222	0.000572
exam_score	0.776751	0.188222	1.000000	0.225713
attendance_percent	-0.031311	0.000572	0.225713	1.000000


A positive correlation can be observed between hours\_studied and exam score, implying that the number of hours studied is associated with the score of exam.

There is a direct relationship between hours studied and exam score.

Creating a correlation heatmap.

```
[65]: x_labels = ["Hours Studied", "Sleep Hours", "Exam Score", "% of Attendance"]
y_labels = ["Exam Score", "Sleep Hours", "Exam Score", "% of Attendance"]
sns.heatmap(corr, annot=True, fmt=".2f", cmap="Blues", xticklabels=x_labels, yticklabels=y_labels)
```


```
[65]: <Axes: >
```



Creating plots showing distributions of study habits and performance.

```
[66]: plot = scores[["hours_studied", "sleep_hours", "previous_scores", "exam_score"]].hist(bins=20)
plt.suptitle("Distribution of Exam Scores and Habits of Students")
```

```
[66]: Text(0.5, 0.98, 'Distribution of Exam Scores and Habits of Students')
```



Splitting numeric data to train and test sets.

```
[39]: y = scores[["exam_score"]]
x = scores.drop(["exam_score", "student_id", "isPassed"], axis=1)

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
random_state=100)
```

A linear regression model was fitted on the training data.

```
[40]: lr = LinearRegression()
lr.fit(x_train, y_train)
```

```
[40]: LinearRegression()
```

Generating predictions for both training and test set.

```
[72]: y_lr_train_pred = lr.predict(x_train)
y_lr_test_pred = lr.predict(x_test)
print(y_lr_train_pred[:5])
```

```
[72]: array([[27.88123664],  
           [36.15845932],  
           [33.77975984],  
           [34.72220155],  
           [22.62250902]])
```

```
[73]: print(y_lr_test_pred[:5])
```

```
[[28.96606215]  
[20.85361165]  
[36.39543205]  
[39.73621892]  
[31.87043094]]
```

Model performance was evaluated using MSE and R2 for both training and test set.

```
[43]: lr_train_mse = mean_squared_error(y_train, y_lr_train_pred)  
lr_train_r2 = r2_score(y_train, y_lr_train_pred)  
  
lr_test_mse = mean_squared_error(y_test, y_lr_test_pred)  
lr_test_r2 = r2_score(y_test, y_lr_test_pred)
```

```
[44]: print("LS MSE (Train): ", lr_train_mse)  
print("LS R2 (Train): ", lr_train_r2)  
print("LS MSE (Test): ", lr_test_mse)  
print("LS R2 (Test): ", lr_test_r2)
```

```
LS MSE (Train):  7.167182016317801  
LS R2 (Train):  0.8412909628396832  
LS MSE (Test):  8.37137635264973  
LS R2 (Test):  0.8205275135283744
```

```
[45]: lr_results = pd.DataFrame(["Linear Regression", lr_train_mse, lr_train_r2,  
                                ↴lr_test_mse, lr_test_r2]).transpose()  
lr_results.columns = ["Method", "Training MSE", "Training R2", "Test MSE",  
                     ↴"Test R2"]  
lr_results
```

	Method	Training MSE	Training R2	Test MSE	Test R2
0	Linear Regression	7.167182	0.841291	8.371376	0.820528

A random forest regressor model was fitted on the training data.

```
[46]: rf = RandomForestRegressor(max_depth=2, random_state=100)  
y_train_1d = y_train.values.ravel()  
rf.fit(x_train, y_train_1d)
```

```
[46]: RandomForestRegressor(max_depth=2, random_state=100)
```

```
[47]: y_rf_train_pred = rf.predict(x_train)
y_rf_test_pred = rf.predict(x_test)
```

Generating predictions for both data sets.

```
[48]: rf_train_mse = mean_squared_error(y_train, y_rf_train_pred)
rf_train_r2 = r2_score(y_train, y_rf_train_pred)

rf_test_mse = mean_squared_error(y_test, y_rf_test_pred)
rf_test_r2 = r2_score(y_test, y_rf_test_pred)
```

```
[49]: rf_results = pd.DataFrame(["Random Forest", rf_train_mse, rf_train_r2,
                                rf_test_mse, rf_test_r2]).transpose()
rf_results.columns = ["Method", "Training MSE", "Training R2", "Test MSE",
                      "Test R2"]
rf_results
```

```
[49]:      Method Training MSE Training R2    Test MSE    Test R2
0  Random Forest     14.004342     0.68989  17.062459  0.634201
```

Comparison of two used models (linear regression and random forest).

```
[50]: df_models = pd.concat([lr_results, rf_results], axis=0).reset_index(drop=True)
df_models
```

```
[50]:      Method Training MSE Training R2    Test MSE    Test R2
0  Linear Regression     7.167182     0.841291  8.371376  0.820528
1      Random Forest     14.004342     0.68989  17.062459  0.634201
```

Visualization of predictions.

```
[51]: plt.figure(figsize=[5,5])
plt.scatter(x=y_train, y=y_lr_train_pred, alpha=0.5)

plt.ylabel("Predict Exam Score")
plt.xlabel("Experimental Exam Score")
plt.plot()
```

```
[51]: []
```

